


Probabilistic Artificial Intelligence

Reasoning over time




2023-2024

1

Approximate Inference

Recap

- Probability: useful for representing beliefs of our agents
 - ▷ Bayes' rule to update beliefs
- Compactly representing
 - ▷ Bayesian networks
 - ▷ exploits conditional independence
- Inference intractable in general...
 - approximate inference



It allows to *maintain* a belief

- Given conditional independent observations
 $P(o_1, o_2 | State) = P(o_1 | State)P(o_2 | State)$


- ...we can also sequentially update:

$$P'(State) := P(o_1 | State)P(State) / P(o_1)$$

$$P''(State) := P(o_2 | State)P'(State) / P(o_2)$$

- ▷ then $P''(State) = P(State | o_1, o_2)$
- ▷ Exercise...?

- “We will see later how to incorporate robot movement over time”
→ **the time has come!**



It allows to *maintain* a belief

- Given conditional independent observations

$$P(o_1, o_2 | State) = P(o_1 | State)P(o_2 | State)$$

- ...we can also sequentially update:

$$P'(State) := P(o_1 | State)P(State) / P(o_1)$$

$$P''(State) := P(o_2 | State)P'(State) / P(o_2)$$

▷ then $P''(State) = P(State | o_1, o_2)$

▷ Exercise...?

with

the updated belief after observing o_1 .

$$P(S | o_1, o_2) = \frac{P(o_1, o_2 | S)P(S)}{\sum_s P(o_1, o_2 | s)P(s)}$$


$$\{\text{conditional independence}\} = \frac{P(o_2 | S)P(o_1 | S)P(S)}{\sum_s P(o_2 | s)P(o_1 | s)P(s)}$$

$$\{\text{multiply with 1}\} = \frac{1/P(o_1)}{1/P(o_1)} \times \frac{P(o_2 | S)P(o_1 | S)P(S)}{\sum_s P(o_2 | s)P(o_1 | s)P(s)}$$

$$\{\text{push inward}\} = \frac{P(o_2 | S) \frac{P(o_1 | S)P(S)}{P(o_1)}}{\sum_s P(o_2 | s) \frac{P(o_1 | s)P(s)}{P(o_1)}}$$

$$= \frac{P(o_2 | S)P'(S)}{\sum_s P(o_2 | s)P'(s)} = P''(\text{state})$$

Warm Up Exercise



- Assume the history is:
(see nothing, move right, see door, move right, see door)
 - ▷ where are we now?
 - ▷ where did we start?

Motivation

Why do we need special methods for time?

- Bayes rule is awesome...
 - ▷ but perhaps not (directly) sufficient to deal with time...?
 - ▷ $P(\text{State} | o_1) := P(o_1 | \text{State})P(\text{State}) / P(o_1)$
- Problems?

Why do we need special methods for time?

■ Bayes rule is awesome...

- ▷ but perhaps not (directly) sufficient to deal with time...?
- ▷ $P(\text{State} | o_1) := P(o_1 | \text{State})P(\text{State}) / P(o_1)$

■ Problems?

E.g.: a self-built drone that determines its position based on images it takes with its camera...

- to make it easier the designers place a landmark
- accuracy is important... the measurement is noisy... so designers decide to average all measurements over the course of 20s...

Will it work?

The World Changes

- ... position could vary substantially over 20s interval

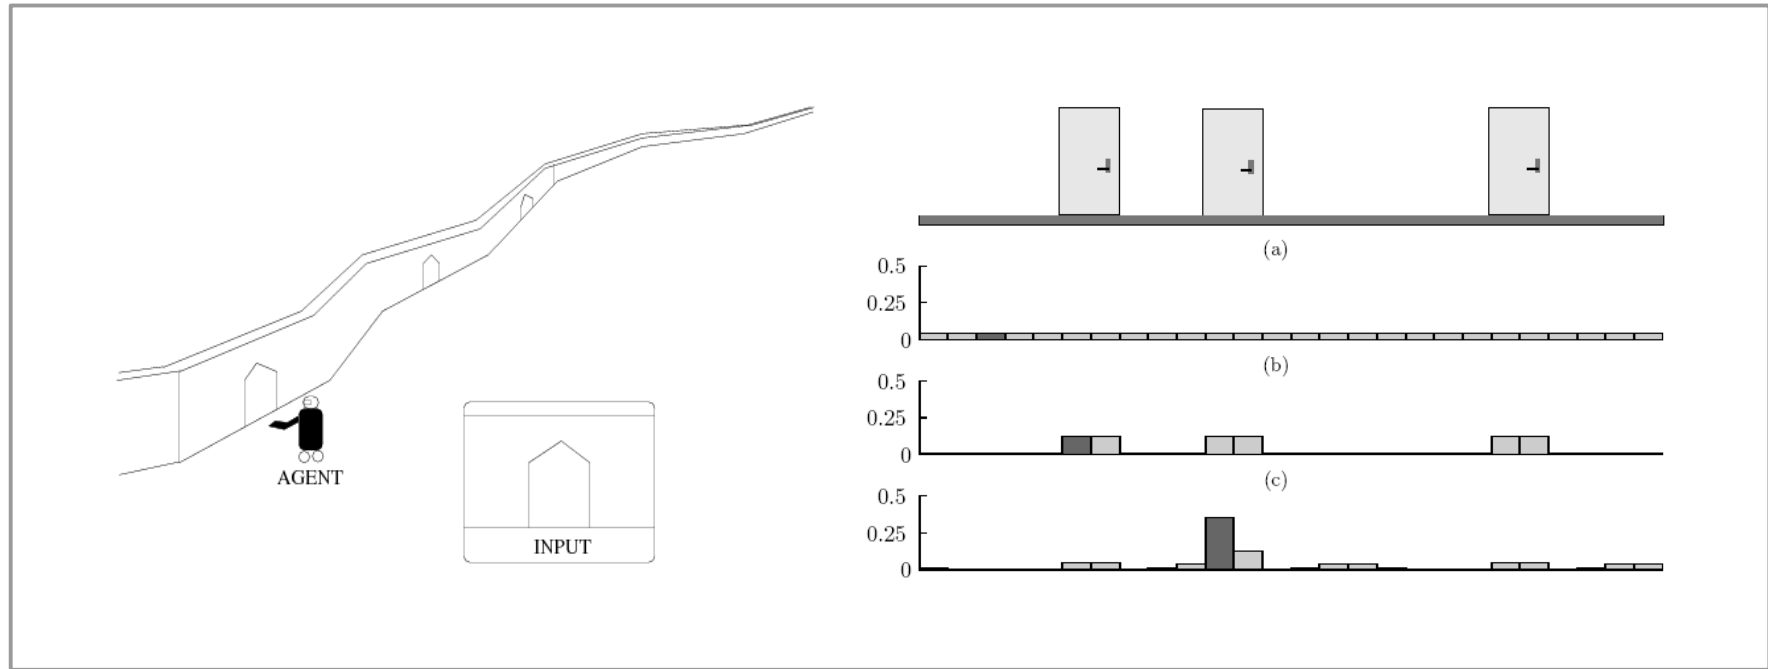
The World Changes

- ... position could vary substantially over 20s interval
- Another example: treating a diabetic patient
 - ▷ need to decide on food intake and insulin dose
 - ▷ need to estimate: current blood sugar and insulin levels – these can vary quickly

The World Changes

- ... position could vary substantially over 20s interval
- Another example: treating a diabetic patient
 - ▷ need to decide on food intake and insulin dose
 - ▷ need to estimate: current blood sugar and insulin levels – these can vary quickly
- Upshot:
 - ▷ need not only consider “sensor model” $P(o_1 | State)$
 - ▷ but also a “transition model” that predicts how the state changes over time.

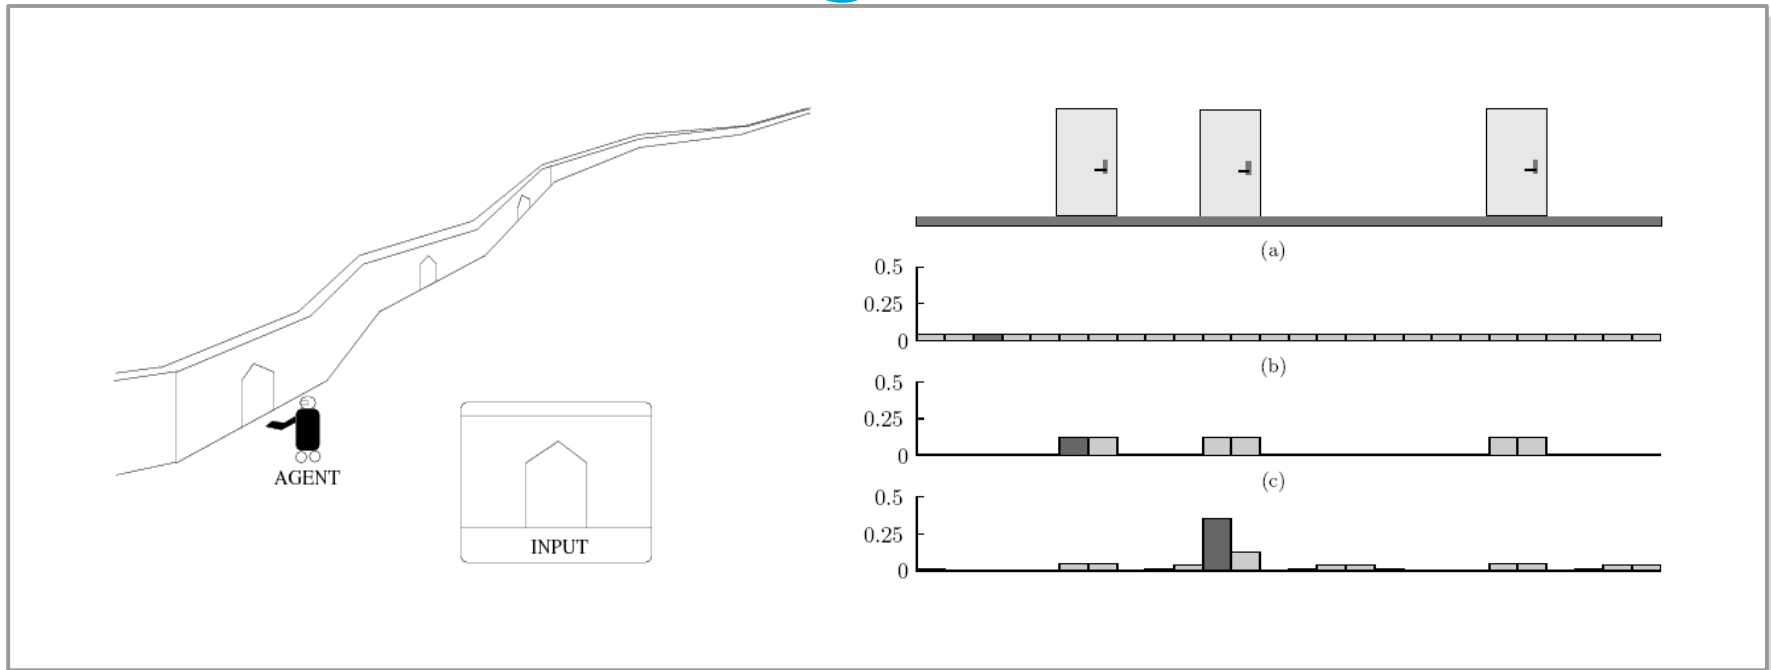
The World Changes



■ Upshot:

- ▷ need not only consider “sensor model” $P(o_1 | State)$
- ▷ but also a “transition model” that predicts how the state changes over time.

The World Changes



■ Upshot:

- ▷ need not only consider “sensor model” $P(o_t | State)$
- ▷ but also a “transition model” that predicts how the state changes over time.

Representing Time

Representing Time

Basic idea: copy state and evidence variables for each time step

\mathbf{X}_t = set of unobservable state variables at time t
e.g., *BloodSugar_t*, *StomachContents_t*, etc.

\mathbf{E}_t = set of observable evidence variables at time t
e.g., *MeasuredBloodSugar_t*, *PulseRate_t*, *FoodEaten_t*

This assumes **discrete time**; step size depends on problem

Notation: $\mathbf{X}_{a:b} = \mathbf{X}_a, \mathbf{X}_{a+1}, \dots, \mathbf{X}_{b-1}, \mathbf{X}_b$

Representing Time

Basic idea: copy state and evidence variables for each time step

\mathbf{X}_t = set of unobservable state variables at time t
e.g., *BloodSugar_t*, *StomachContents_t*, etc.

\mathbf{E}_t = set of observable evidence variables at time t
e.g., *MeasuredBloodSugar_t*, *PulseRate_t*, *FoodEaten_t*

This assumes **discrete time**; step size depends on problem

Notation: $\mathbf{X}_{a:b} = \mathbf{X}_a, \mathbf{X}_{a+1}, \dots, \mathbf{X}_{b-1}, \mathbf{X}_b$

Question:

How can we compactly represent a probability distribution over the evolution of state \mathbf{X}_t ?

Markov Processes (Markov Chains)

Construct a Bayes net from these variables: parents?

Markov Processes (Markov Chains)

Construct a Bayes net from these variables: parents?

Markov assumption: \mathbf{X}_t depends on **bounded** subset of $\mathbf{X}_{0:t-1}$

First-order Markov process: $\mathbf{P}(\mathbf{X}_t | \mathbf{X}_{0:t-1}) = \mathbf{P}(\mathbf{X}_t | \mathbf{X}_{t-1})$

..... 'transition model'

First-order



Markov Processes (Markov Chains)


Construct a Bayes net from these variables: parents?

Markov assumption: \mathbf{X}_t depends on **bounded** subset of $\mathbf{X}_{0:t-1}$

First-order Markov process: $\mathbf{P}(\mathbf{X}_t | \mathbf{X}_{0:t-1}) = \mathbf{P}(\mathbf{X}_t | \mathbf{X}_{t-1})$

Second-order Markov process: $\mathbf{P}(\mathbf{X}_t | \mathbf{X}_{0:t-1}) = \mathbf{P}(\mathbf{X}_t | \mathbf{X}_{t-2}, \mathbf{X}_{t-1})$

.....
'transition
model'



Markov Processes (Markov Chains)


Construct a Bayes net from these variables: parents?

Markov assumption: \mathbf{X}_t depends on **bounded** subset of $\mathbf{X}_{0:t-1}$

First-order Markov process: $\mathbf{P}(\mathbf{X}_t | \mathbf{X}_{0:t-1}) = \mathbf{P}(\mathbf{X}_t | \mathbf{X}_{t-1})$

Second-order Markov process: $\mathbf{P}(\mathbf{X}_t | \mathbf{X}_{0:t-1}) = \mathbf{P}(\mathbf{X}_t | \mathbf{X}_{t-2}, \mathbf{X}_{t-1})$

.....
'transition
model'

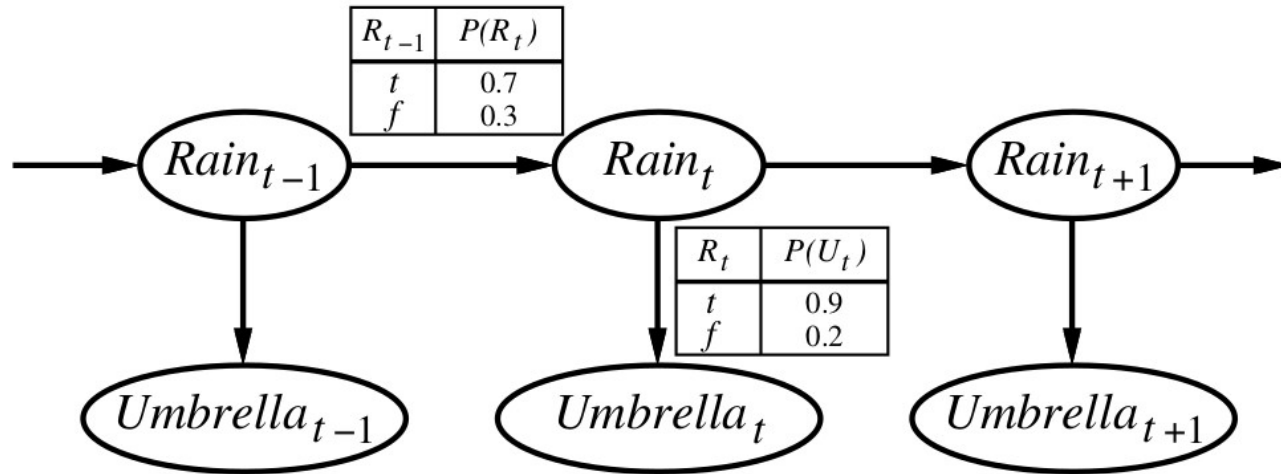


Sensor Markov assumption: $\mathbf{P}(\mathbf{E}_t | \mathbf{X}_{0:t}, \mathbf{E}_{0:t-1}) = \mathbf{P}(\mathbf{E}_t | \mathbf{X}_t)$

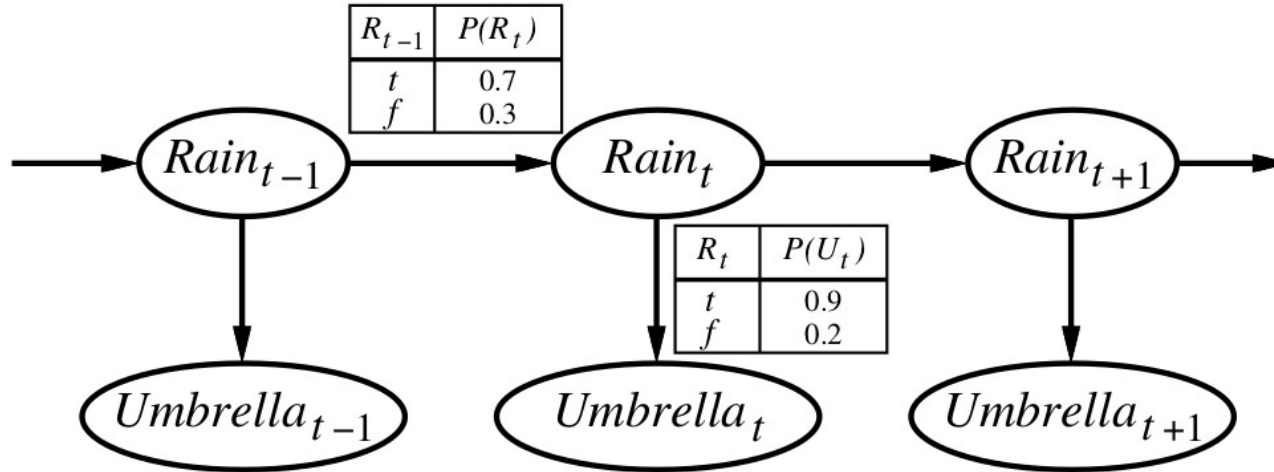
.....
'observation
model'

Stationary process: transition model $\mathbf{P}(\mathbf{X}_t | \mathbf{X}_{t-1})$ and
sensor model $\mathbf{P}(\mathbf{E}_t | \mathbf{X}_t)$ fixed for all t

Example



Example



This is called a **hidden Markov model (HMM)**

- state consists of a single discrete variable

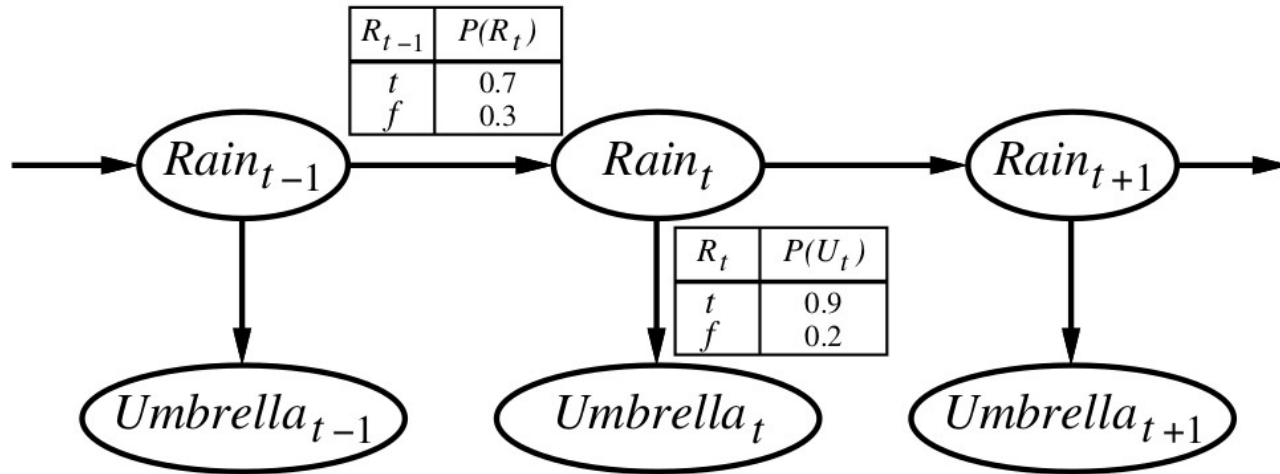
Semantics of Markov Chains

- Transition model + observation model
+ initial state distribution $P(\mathbf{X}_0)$ = joint PD

$$P(\mathbf{X}_{0:t}, \mathbf{E}_{1:t}) = P(\mathbf{X}_0) \prod_{i=1:t} P(\mathbf{X}_i | \mathbf{X}_{i-1}) P(\mathbf{E}_i | \mathbf{X}_i)$$


- “Unrolled over time” it is just a Bayesian network
 - ▷ note: we typically assume that $P(\mathbf{X}_0)$ captures all our knowledge at $t=0$, and hence there is no observation \mathbf{E}_0 (as a convention)

The Markov Assumption...?



Is it appropriate?

The Markov Assumption...?



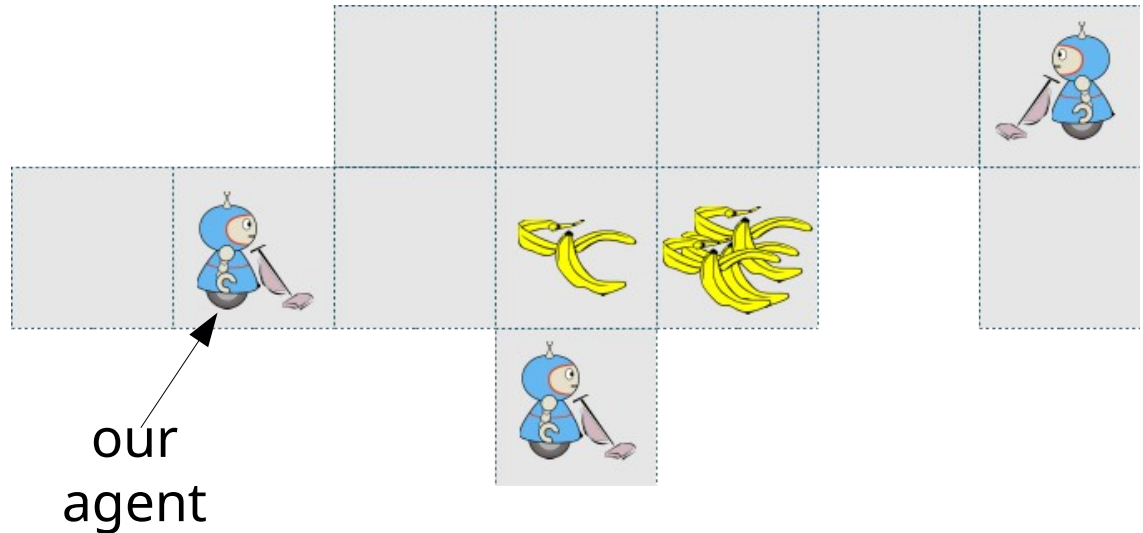
First-order Markov assumption not exactly true in real world!

Possible fixes:

1. **Increase order** of Markov process
2. **Augment state**, e.g., add $Temp_t$, $Pressure_t$

Markov Assumption - 2


- Or how about “Spatial Task Allocation Problems”?
 - ▷ Need to decide where we go to clean...
 - ▷ How will the state change?
 - ▷ What information do we need?




Inference in Hidden Markov Models

Inference on Markov Chains

- Unrolled we have a BN...
 - ▷ So we can use VE, rejection sampling, likelihood weighting etc...
- ...but what can we ask?




Types of Queries



Filtering: $P(\mathbf{X}_t | \mathbf{e}_{1:t})$

belief state—input to the decision process of a rational agent

Types of Queries




Filtering: $P(\mathbf{X}_t | \mathbf{e}_{1:t})$

belief state—input to the decision process of a rational agent

Prediction: $P(\mathbf{X}_{t+k} | \mathbf{e}_{1:t})$ for $k > 0$

evaluation of possible action sequences;
like filtering without the evidence

Types of Queries



Filtering: $P(\mathbf{X}_t | \mathbf{e}_{1:t})$

belief state—input to the decision process of a rational agent


Prediction: $P(\mathbf{X}_{t+k} | \mathbf{e}_{1:t})$ for $k > 0$

evaluation of possible action sequences;
like filtering without the evidence

Smoothing: $P(\mathbf{X}_k | \mathbf{e}_{1:t})$ for $0 \leq k < t$

better estimate of past states, essential for learning

Types of Queries



Filtering: $P(\mathbf{X}_t | \mathbf{e}_{1:t})$

belief state—input to the decision process of a rational agent

Prediction: $P(\mathbf{X}_{t+k} | \mathbf{e}_{1:t})$ for $k > 0$

evaluation of possible action sequences;
like filtering without the evidence


Smoothing: $P(\mathbf{X}_k | \mathbf{e}_{1:t})$ for $0 \leq k < t$

better estimate of past states, essential for learning

Most likely explanation: $\arg \max_{\mathbf{x}_{1:t}} P(\mathbf{x}_{1:t} | \mathbf{e}_{1:t})$

speech recognition, decoding with a noisy channel

Types of Queries



Filtering: $P(\mathbf{X}_t | \mathbf{e}_{1:t})$

belief state—input to the decision process of a rational agent

Prediction: $P(\mathbf{X}_{t+k} | \mathbf{e}_{1:t})$ for $k > 0$

evaluation of possible action sequences;
like filtering without the evidence

Smoothing: $P(\mathbf{X}_k | \mathbf{e}_{1:t})$ for $0 \leq k < t$

better estimate of past states, essential for learning


Most likely explanation: $\arg \max_{\mathbf{x}_{1:t}} P(\mathbf{x}_{1:t} | \mathbf{e}_{1:t})$

speech recognition, decoding with a noisy channel

many names:
belief tracking, state
estimation, recursive
Bayesian estimation,
etc.

Focus: Inference in HMMs

- We focus on HMMs, but...
 - ▷ ...can be generalized (cf. R&N)
 - ▷ e.g., continuous states: Kalman filter




- **Inference in HMMs:** answering queries *given* the HMM parameters
 - ▷ If we are trying to infer those parameters themselves: **learning HMMs**

Filtering: $b_t(X_t) = P(X_t | e_{1:t})$

■ Ideal: a recursive way to compute

- ▷ new belief $b_{t+1}(X_{t+1})$
- ▷ from old belief $b_t(X_t)$




the term 'belief' and notation $b(X)$ (or $b(s)$) is very common in reinforcement learning and planning under uncertainty

Filtering: $\mathbf{b}_t(X_t) = P(X_t | e_{1:t})$

- Ideal: a recursive way to compute

- ▷ new belief $\mathbf{b}_{t+1}(X_{t+1})$
- ▷ from old belief $\mathbf{b}_t(X_t)$




- Why?

Filtering: $\mathbf{b}_t(X_t) = P(X_t | e_{1:t})$

- Ideal: a recursive way to compute

- ▷ new belief $\mathbf{b}_{t+1}(X_{t+1})$
- ▷ from old belief $\mathbf{b}_t(X_t)$




- How?

Filtering: $\mathbf{b}_t(X_t) = P(X_t | e_{1:t})$

■ Ideal: a recursive way to compute

- ▷ new belief $\mathbf{b}_{t+1}(X_{t+1})$
- ▷ from old belief $\mathbf{b}_t(X_t)$




$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1} | e_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) \\ &= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned}$$

Filtering: $\mathbf{b}_t(X_t) = P(X_t | e_{1:t})$

■ Ideal: a recursive way to compute

- ▷ new belief $\mathbf{b}_{t+1}(X_{t+1})$
- ▷ from old belief $\mathbf{b}_t(X_t)$




$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1} | e_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) \quad ? \\ &= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned}$$

Filtering: $b_t(X_t) = P(X_t | e_{1:t})$

■ Ideal: a recursive way to compute

- ▷ new belief $b_{t+1}(X_{t+1})$
- ▷ from old belief $b_t(X_t)$




$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1} | e_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) \\ &= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned}$$

$$\begin{aligned} P(X_{t+1} | e_{1:t}, e_{t+1}) &= \\ P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) / P(e_{t+1} | e_{1:t}) \end{aligned}$$

Filtering: $b_t(X_t) = P(X_t | e_{1:t})$

■ Ideal: a recursive way to compute

- ▷ new belief $b_{t+1}(X_{t+1})$
- ▷ from old belief $b_t(X_t)$




$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1} | e_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) \\ &= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned} \quad \begin{array}{c} \curvearrowright \\ ? \end{array}$$

Filtering: $b_t(X_t) = P(X_t | e_{1:t})$

■ Ideal: a recursive way to compute

- ▷ new belief $b_{t+1}(X_{t+1})$
- ▷ from old belief $b_t(X_t)$




$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1} | e_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) \\ &= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned}$$

sensor Markov assumption!

Filtering: $\mathbf{b}_t(X_t) = P(X_t | e_{1:t})$

■ Ideal: a recursive way to compute

- ▷ new belief $\mathbf{b}_{t+1}(X_{t+1})$
- ▷ from old belief $\mathbf{b}_t(X_t)$




$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1} | e_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) \\ &= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned}$$

I.e., prediction + estimation.

Filtering: $\mathbf{b}_t(X_t) = P(X_t | e_{1:t})$

■ Ideal: a recursive way to compute

- ▷ new belief $\mathbf{b}_{t+1}(X_{t+1})$
- ▷ from old belief $\mathbf{b}_t(X_t)$




$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1} | e_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) \\ &= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned}$$

I.e., prediction + estimation.

Filtering: $b_t(X_t) = P(X_t | e_{1:t})$

Prediction:




$$= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t})$$

I.e., prediction + estimation.

Filtering: $b_t(X_t) = P(X_t | e_{1:t})$

and estimation:




$$= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t})$$

I.e., prediction + estimation.

Filtering: $\mathbf{b}_t(X_t) = P(X_t | e_{1:t})$

■ Ideal: a recursive way to compute

- ▷ new belief $\mathbf{b}_{t+1}(X_{t+1})$
- ▷ from old belief $\mathbf{b}_t(X_t)$



$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1} | e_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) \\ &= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned}$$

I.e., **prediction** + **estimation**. Prediction by summing out X_t :


$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= \alpha P(e_{t+1} | X_{t+1}) \sum_{\mathbf{x}_t} P(X_{t+1} | \mathbf{x}_t, e_{1:t}) P(\mathbf{x}_t | e_{1:t}) \\ &= \alpha P(e_{t+1} | X_{t+1}) \sum_{\mathbf{x}_t} P(X_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | e_{1:t}) \end{aligned}$$

↘ ?

Filtering: $\mathbf{b}_t(X_t) = P(X_t | e_{1:t})$

■ Ideal: a recursive way to compute

- ▷ new belief $\mathbf{b}_{t+1}(X_{t+1})$
- ▷ from old belief $\mathbf{b}_t(X_t)$



$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= P(X_{t+1} | e_{1:t}, e_{t+1}) \\ &= \alpha P(e_{t+1} | X_{t+1}, e_{1:t}) P(X_{t+1} | e_{1:t}) \\ &= \alpha P(e_{t+1} | X_{t+1}) P(X_{t+1} | e_{1:t}) \end{aligned}$$


I.e., **prediction** + **estimation**. Prediction by summing out X_t :

$$\begin{aligned} P(X_{t+1} | e_{1:t+1}) &= \alpha P(e_{t+1} | X_{t+1}) \sum_{\mathbf{x}_t} P(X_{t+1} | \mathbf{x}_t, e_{1:t}) P(\mathbf{x}_t | e_{1:t}) \\ &= \alpha P(e_{t+1} | X_{t+1}) \sum_{\mathbf{x}_t} P(X_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | e_{1:t}) \end{aligned}$$

Markov
assumption

Filtering: $\mathbf{b}_t(X_t) = P(X_t | e_{1:t})$

- Ideal: a recursive way to compute
 - ▷ new belief $\mathbf{b}_{t+1}(X_{t+1})$



So now, we have this expression of the form

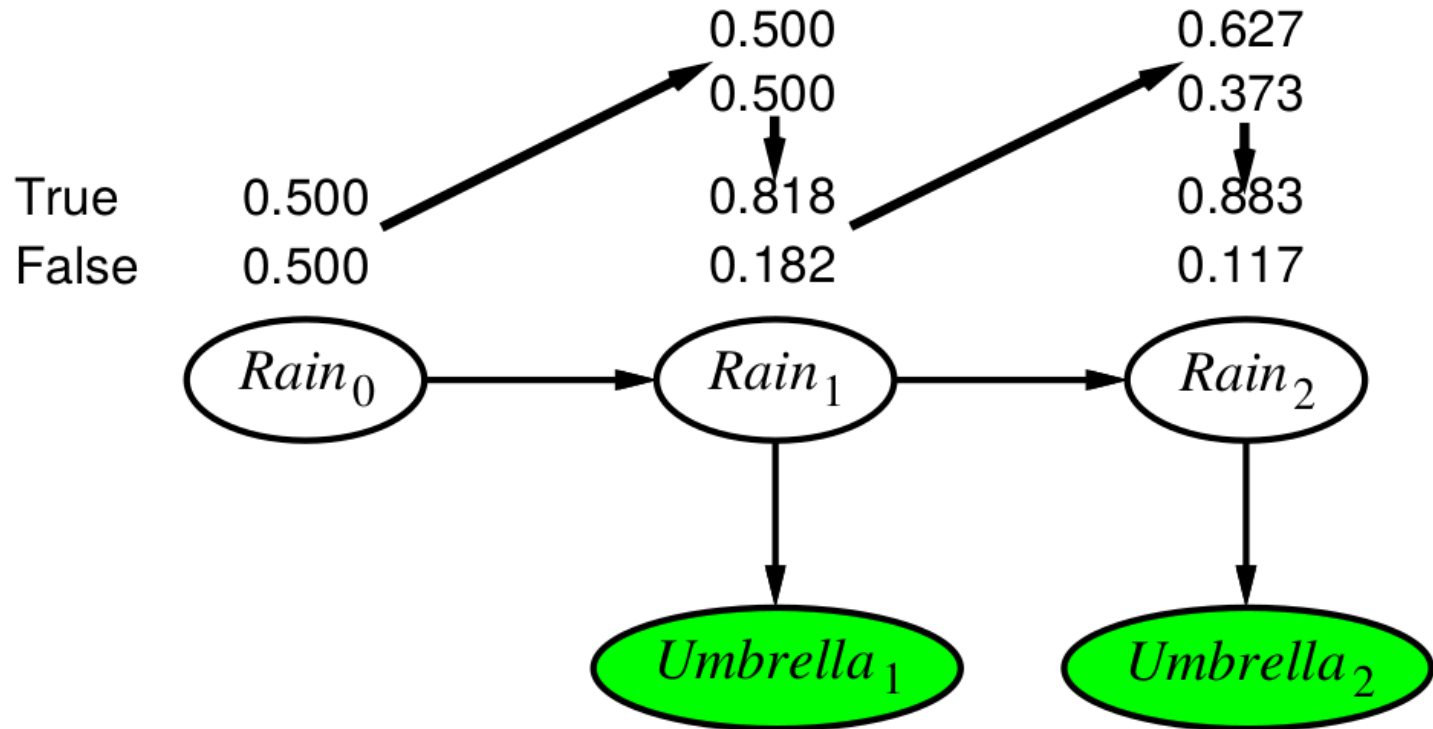
$$\begin{aligned}\mathbf{b}_{t+1}(X_{t+1}) &= \alpha P(e_{t+1} | X_{t+1}) \sum_x P(X_{t+1} | x_t) \mathbf{b}_t(x_t) \\ &= \alpha \text{Forward}(\mathbf{b}_t(X_t), e_{t+1})\end{aligned}$$

- the $\mathbf{b}_t(X_t)$ are also called **forward messages**, $\mathbf{f}_{1:t}$
- initialize: $\mathbf{f}_{1:0} = P(X_0)$

$$= \alpha P(\mathbf{e}_{t+1} | \mathbf{X}_{t+1}) \sum_{\mathbf{x}_t} P(\mathbf{X}_{t+1} | \mathbf{x}_t) P(\mathbf{x}_t | \mathbf{e}_{1:t})$$


Filtering Example

Numbers?
→ Exercise!



Prediction

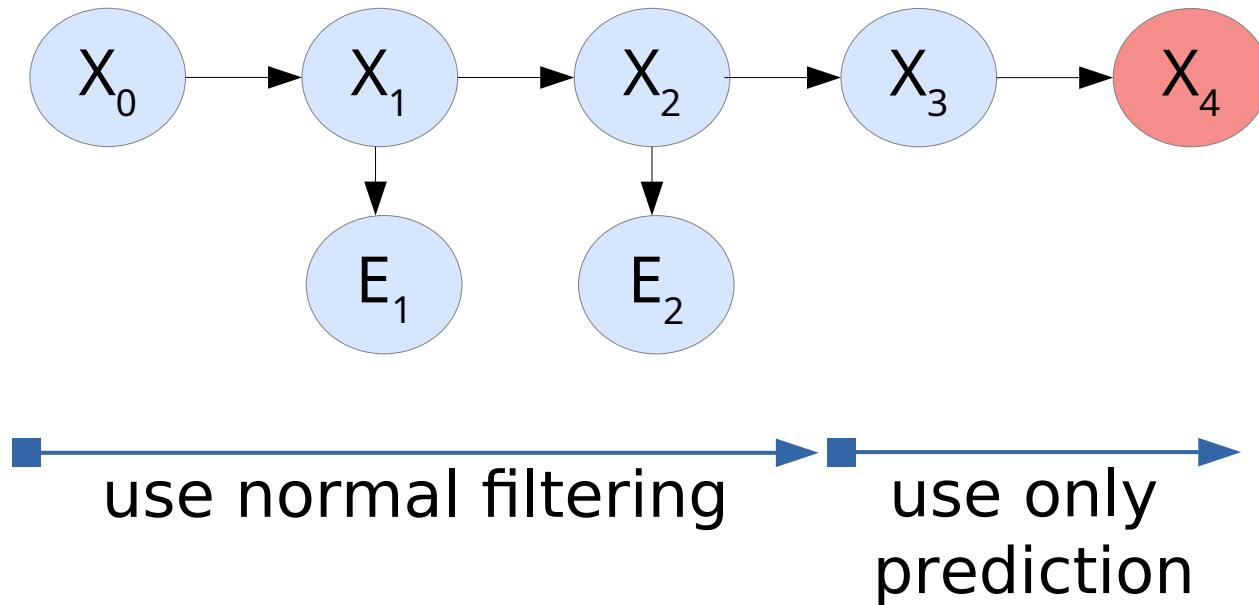
■ Query: $P(X_{t+K} \mid e_{1:t})$?



Prediction

$$\mathbf{b}_{t+1}(X_{t+1}) = \alpha \mathbf{P}(e_{t+1} | X_{t+1}) \sum_x \mathbf{P}(X_{t+1} | x_t) \mathbf{b}_t(x_t)$$

■ Query: $\mathbf{P}(X_{t+K} | e_{1:t})$?



Prediction... limits?


- How well can we predict “Rain” 10 steps into the future...?

Prediction... limits?

- How well can we predict “Rain” 10 steps into the future...?
- Belief will converge to $\langle 0.5, 0.5 \rangle$ quite fast...
 - ▷ called the **stationary distribution**
 - ▷ the more stochastic the process, the shorter the **mixing time**
 - ▷ predicting beyond a fraction of the mixing time will not work


'Smoothing'

- Query: $P(X_t | e_{1:t+K})$




- ▷ more information available than with filtering → less variance in answer → 'smoother result'
- Now: also need to take into account information (passed back) from the future... **forward-backward algorithm**

Smoothing: Approach



$$\begin{aligned} \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:t}) &= \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &= \alpha \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}) \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{e}_{1:k}) \quad \curvearrowright \quad ? \\ &= \alpha \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}) \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) \\ &= \alpha \mathbf{f}_{1:k} \mathbf{b}_{k+1:t} \end{aligned}$$


Smoothing: Approach



$$\begin{aligned} \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:t}) &= \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}, \mathbf{e}_{k+1:t}) \\ &= \alpha \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}) \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{e}_{1:k}) \\ &= \alpha \mathbf{P}(\mathbf{X}_k | \mathbf{e}_{1:k}) \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) \\ &= \alpha \mathbf{f}_{1:k} \mathbf{b}_{k+1:t} \end{aligned}$$


product of
forward
and
backward
messages

Smoothing: Backward messages




$$\mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) = \sum_{\mathbf{x}_{k+1}} \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k)$$

Smoothing: Backward messages




$$\begin{aligned} \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) &= \sum_{\mathbf{x}_{k+1}} \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\ &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \end{aligned}$$

Smoothing: Backward messages



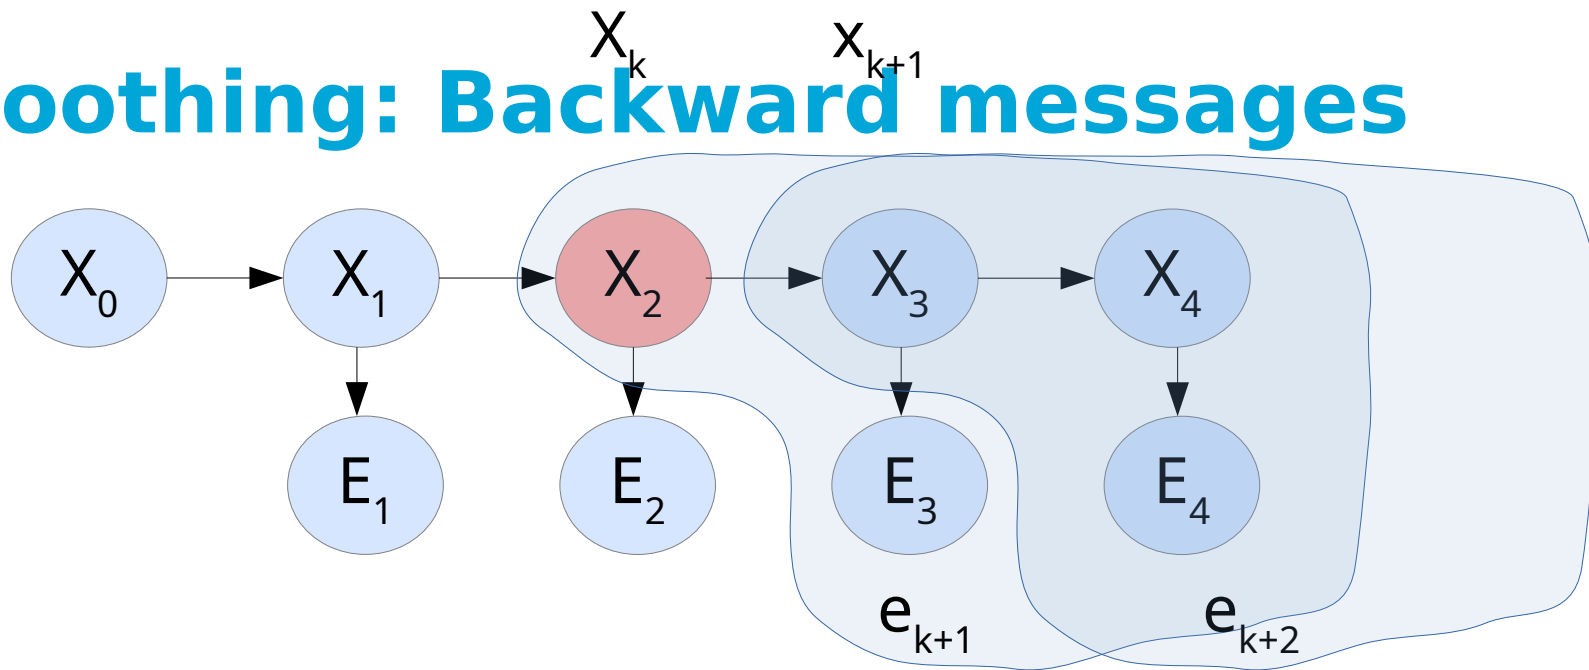
$$\begin{aligned} \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) &= \sum_{\mathbf{x}_{k+1}} \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\ &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\ &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1} | \mathbf{x}_{k+1}) P(\mathbf{e}_{k+2:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \end{aligned}$$

Smoothing: Backward messages




$$\begin{aligned} \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) &= \sum_{\mathbf{x}_{k+1}} \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\ &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\ &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1} | \mathbf{x}_{k+1}) P(\mathbf{e}_{k+2:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \end{aligned}$$

Smoothing: Backward messages



$$\begin{aligned}
 \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) &= \sum_{\mathbf{x}_{k+1}} \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\
 &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\
 &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1} | \mathbf{x}_{k+1}) P(\mathbf{e}_{k+2:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k)
 \end{aligned}$$

Smoothing: Backward messages

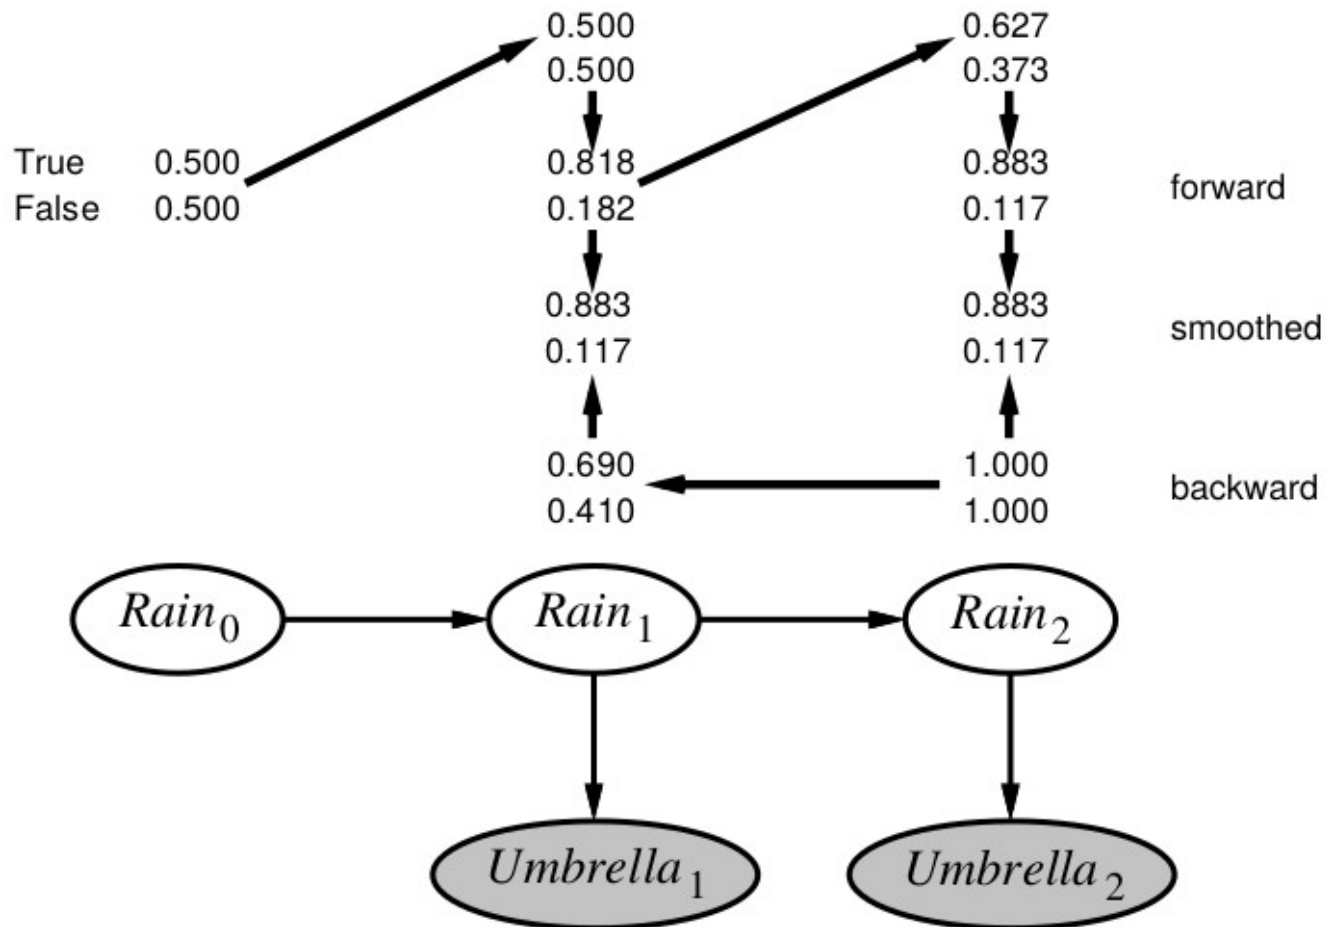


$$\begin{aligned}
 \mathbf{b}_{k+1:t}(X_k) &= \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k) = \sum_{\mathbf{x}_{k+1}} \mathbf{P}(\mathbf{e}_{k+1:t} | \mathbf{X}_k, \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\
 &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k) \\
 &= \sum_{\mathbf{x}_{k+1}} P(\mathbf{e}_{k+1} | \mathbf{x}_{k+1}) P(\mathbf{e}_{k+2:t} | \mathbf{x}_{k+1}) \mathbf{P}(\mathbf{x}_{k+1} | \mathbf{X}_k)
 \end{aligned}$$

initialize: $\mathbf{b}_{t+1:t} = \mathbf{1}$ (i.e., vector)

$\mathbf{b}_{k+2:t}(\mathbf{x}_{k+1})$

Forward-Backward Illustrated



Smoothing vs most-likely sequences

- Smoothing can compute

$$\{P(X_0 | E_{1:t}), P(X_1 | E_{1:t}), \dots, P(X_t | E_{1:t})\}$$

- But does not give most likely sequence!

$$\max_{(x_0, x_1, \dots, x_t)} P(x_0, x_1, \dots, x_t | E_{1:t})$$

!=

$$\{ \max_{x_0} P(x_0 | E_{1:t}), \max_{x_1} P(x_1 | E_{1:t}), \dots, \max_{x_t} P(x_t | E_{1:t}) \}$$

- (Need a different algo: Viterby – see book)

Dynamic Bayesian Networks

Complex worlds

BOUTILIER, DEAN, & HANKS

- E.g., the world has many aspects





Figure 1: A decision-theoretic planning problem

Complex worlds

BOUTILIER, DEAN, & HANKS

- E.g., the world has many aspects

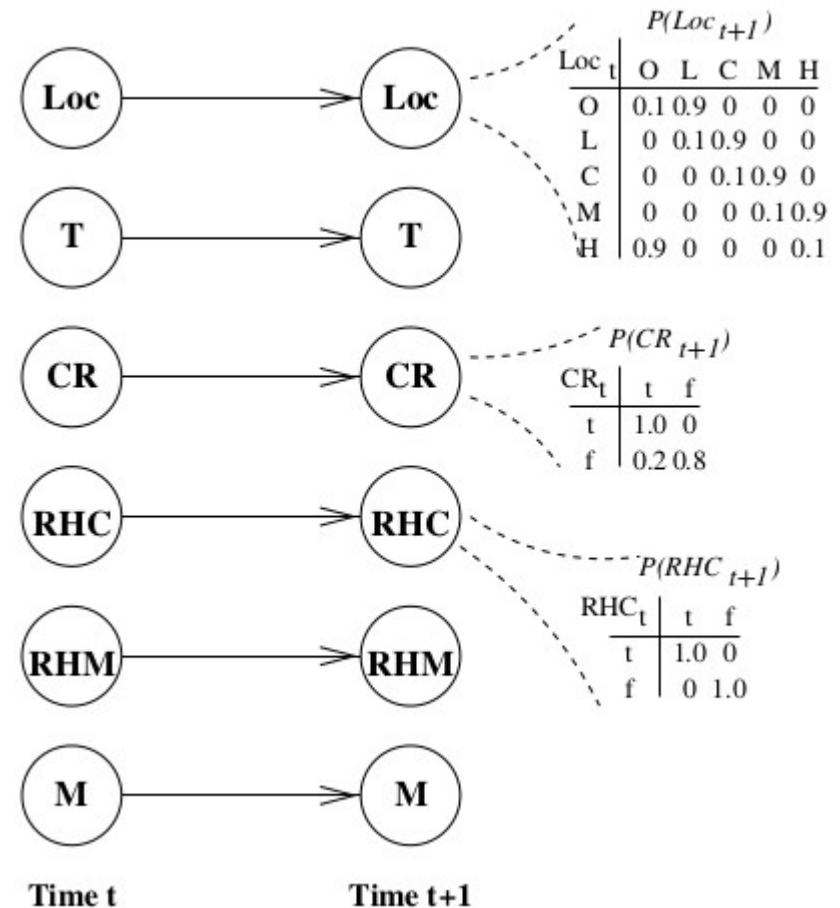


Features	Denoted	Description
Location	$Loc(M)$, etc.	Location of robot. Five possible locations: mailroom (M), coffee room (C), user's office (O), hallway (H), laboratory (L)
Tidiness	$T(0)$, etc.	Degree of lab tidiness. Five possible values: from 0 (messiest) to 4 (tidiest)
Mail present	M, \overline{M}	Is there mail in user's mail box? True (M) or False (\overline{M})
Robot has mail	RHM, \overline{RHM}	Does the robot have mail in its possession?
Coffee request	CR, \overline{CR}	Is there an outstanding (unfulfilled) request for coffee by the user?
Robot has coffee	RHC, \overline{RHC}	Does the robot have coffee in its possession?

m





Again: compact representations

- E.g.
“move counter-clockwise”



Again: compact representations

Every HMM is a single-variable DBN; every discrete DBN is an HMM




Sparse dependencies \Rightarrow exponentially fewer parameters;


e.g., 20 state variables, three parents each

DBN has $20 \times 2^3 = 160$ parameters, HMM has $2^{20} \times 2^{20} \approx 10^{12}$


Example, dealing with failing sensors




- Can explicitly represent prob. of sensors failing



Example, dealing with failing sensors



- battery meter was '5' for 20 steps
- complete discharge unlikely (according to transition model)
- meter can have fluke
- meter might be broken




Inference in DBNs...

Suggestions?

Difficulties?

Inference in DBNs...


Naive method: **unroll** the network and run any exact algorithm



Problem: inference cost for each update grows with t

Inference in DBNs...

Naive method: **unroll** the network and run any exact algorithm



Problem: inference cost for each update grows with t


- But filtering can be done recursively...

$$\mathbf{b}_{t+1}(X_{t+1}) = \alpha \mathbf{P}(e_{t+1} | X_{t+1}) \sum_x \mathbf{P}(X_{t+1} | x_t) \mathbf{b}_t(x_t)$$

... so perhaps can run variable elimination per time step?

Inference in DBNs

Naive method: **unroll** the network



Problem: inference cost for each u

Yes... but...

- for a DBN, the factors that you will construct will be huge.
- will include all variables that have parents in previous stage...

→ **approximate inference**

- But filtering can be done recursively...

$$\mathbf{b}_{t+1}(X_{t+1}) = \alpha \mathbf{P}(e_{t+1} | X_{t+1}) \sum_x \mathbf{P}(X_{t+1} | x_t) \mathbf{b}_t(x_t)$$

... so perhaps can run variable elimination per time step?

"entanglement"

→ try this!

- draw a simple DBN with 3 variables.
- try and compute b' assuming that b is completely factored:

$$b(X_1, X_2, X_3) = b(X_1)b(X_2)b(X_3)$$

Ns

Yes... but...

- for a DBN, the factors that you will construct will be huge.
- will include all variables that have parents in previous stage...

→ **approximate inference**

recursively...

$$b_{t+1}(X_{t+1}) = \alpha P(e_{t+1} | X_{t+1}) \sum_x P(X_{t+1} | x_t) b_t(x_t)$$

... so perhaps can run variable elimination per time step?


Particle Filters

Approximate inference for DBNs

- Why not Likelihood Weighting?


- For $P(X_2 \mid \mathbf{e}_{1:2})$ this would be:

- ▷ For $i=1:\text{num_samples}$
 - sample states x_0 , x_1 , and x_2 (from transition model)
 - form data point $(x_0, x_1, x_2, \mathbf{e}_{1:2})$
 - compute 'weight' w
 - set $w[x_2] += w$
- ▷ renormalize weights
- ▷ $P(x_2 \mid \mathbf{e}_{1:2}) = W[x_2]$




LW... Problems for DBNs

- Two main problems....



LW... Problems for DBNs




- Two main problems....

- 1) running each sample from step 1 to t:
time needed grows over time...

- 2)

LW... Problems for DBNs





■ Two main problems....

1) running each sample from step 1 to t :
time needed grows over time...

2) states are sampled **independently** of the evidence

- most samples are completely wrong
 - very few data points $(x_0, x_1, x_2, \mathbf{e})$ will be likely
 - get all the (renormalized) weight

LW... Problems for DBNs





o t:

r of the evidence

$\zeta_0, x_1, x_2, \mathbf{e}$ will be likely

Fixing these: Particle Filtering

run all N samples at the same time
→ 'particles' themselves represent belief



■ Two main problems....

1) running each sample from step 1 to t :
time needed grows over time...

2) states are sampled **independently** of the evidence


- most samples are completely wrong
→ very few data points $(x_0, x_1, x_2, \mathbf{e})$ will be likely
→ get all the (renormalized) weight

resampling
→ focus attention on parts of state space
with large prob. under the evidence

Particle Filtering: Intuition

Basic idea: ensure that the population of samples (“particles”) tracks the high-likelihood regions of the state-space

Replicate particles proportional to likelihood for e_t



Particle Filtering: Updates

Propagate forward: populations of \mathbf{x}_{t+1} are

$$N(\mathbf{x}_{t+1}|\mathbf{e}_{1:t}) = \sum_{\mathbf{x}_t} P(\mathbf{x}_{t+1}|\mathbf{x}_t)N(\mathbf{x}_t|\mathbf{e}_{1:t})$$

Weight samples by their likelihood for \mathbf{e}_{t+1} :

$$W(\mathbf{x}_{t+1}|\mathbf{e}_{1:t+1}) = P(\mathbf{e}_{t+1}|\mathbf{x}_{t+1})N(\mathbf{x}_{t+1}|\mathbf{e}_{1:t})$$

Resample to obtain populations proportional to W :

$$\begin{aligned} N(\mathbf{x}_{t+1}|\mathbf{e}_{1:t+1})/N &= \alpha W(\mathbf{x}_{t+1}|\mathbf{e}_{1:t+1}) = \alpha P(\mathbf{e}_{t+1}|\mathbf{x}_{t+1})N(\mathbf{x}_{t+1}|\mathbf{e}_{1:t}) \\ &= \alpha P(\mathbf{e}_{t+1}|\mathbf{x}_{t+1}) \sum_{\mathbf{x}_t} P(\mathbf{x}_{t+1}|\mathbf{x}_t)N(\mathbf{x}_t|\mathbf{e}_{1:t}) \\ &= \alpha' P(\mathbf{e}_{t+1}|\mathbf{x}_{t+1}) \sum_{\mathbf{x}_t} P(\mathbf{x}_{t+1}|\mathbf{x}_t)P(\mathbf{x}_t|\mathbf{e}_{1:t}) \\ &= P(\mathbf{x}_{t+1}|\mathbf{e}_{1:t+1}) \end{aligned}$$

Particle Filtering: Updates

Propagate forward: populations of \mathbf{x}_{t+1} are

$$N(\mathbf{x}_{t+1}|\mathbf{e}_{1:t}) = \sum_{\mathbf{x}_t} P(\mathbf{x}_{t+1}|\mathbf{x}_t)N(\mathbf{x}_t|\mathbf{e}_{1:t})$$

Weight samples by their likelihood for \mathbf{e}_{t+1} :

$$W(\mathbf{x}_{t+1}|\mathbf{e}_{1:t+1}) = P(\mathbf{e}_{t+1}|\mathbf{x}_{t+1})N(\mathbf{x}_{t+1}|\mathbf{e}_{1:t})$$

Resample to obtain populations proportional to W :

$$\begin{aligned} N(\mathbf{x}_{t+1}|\mathbf{e}_{1:t+1})/N &= \alpha W(\mathbf{x}_{t+1}|\mathbf{e}_{1:t+1}) = \alpha P(\mathbf{e}_{t+1}|\mathbf{x}_{t+1})N(\mathbf{x}_{t+1}|\mathbf{e}_{1:t}) \\ &= \alpha P(\mathbf{e}_{t+1}|\mathbf{x}_{t+1}) \sum_{\mathbf{x}_t} P(\mathbf{x}_{t+1}|\mathbf{x}_t)N(\mathbf{x}_t|\mathbf{e}_{1:t}) \\ &= \alpha' P(\mathbf{e}_{t+1}|\mathbf{x}_{t+1}) \sum_{\mathbf{x}_t} P(\mathbf{x}_{t+1}|\mathbf{x}_t)P(\mathbf{x}_t|\mathbf{e}_{1:t}) \\ &= P(\mathbf{x}_{t+1}|\mathbf{e}_{1:t+1}) \end{aligned}$$


assumption:

consistent at stage t

$$N(\mathbf{x}_t|\mathbf{e}_{1:t})/N = P(\mathbf{x}_t|\mathbf{e}_{1:t})$$

Particle Filtering in Practice

- In practice this works very well
- Many video explanations -E.g.,
 - ▷ <https://www.youtube.com/watch?v=sz7cJuMgKFg>
 - ▷ <https://youtu.be/eAqAFSrTGGY?t=1117>
- Many real world applications
 - ▷ <http://stanford.edu/~cpiech/cs221/apps/driverlessCar.html>



Reasoning over Time: Summary

- Agents need to reason over time: time-slice based Bayesian networks
 - ▷ Hidden Markov Models – HMMs – 1 discrete state
 - ▷ Dynamic Bayesian networks – DBNs
- Inference over time
 - ▷ Filtering, prediction, smoothing
 - ▷ Scaling to large DBNs: intractable
 - approximate inference: particle filter