

Inverse Reinforcement Learning

CS4375 Artificial Intelligence Techniques

Luciano Cavalcante Siebert

Assistant Professor

Interactive Intelligence (II) Group, INSY, EEMCS, TU Delft

l.cavalcantesiebert@tudelft.nl

What we are going to talk about...

- Introduction to Imitation learning
- Inverse Reinforcement Learning (IRL)
 - Problem definition
 - Linear formulation and example
 - Other approaches: Max Entropy IRL, Adversarial IRL, Collaborative IRL
 - Limitations and critiques

Introduction

For computer games, the reward is usually quite clear:



However, in real world applications this is often not the case. Often a proxy is used as reward:



Imitation learning


- Instead of trying to learn from sparse rewards or manually specified reward function, an expert provides a set of demonstrations
- Agent then tries to learn from these demonstrations
- Approaches to imitation learning include:
 - Behavioral cloning
 - Direct policy learning
 - **Inverse reinforcement learning (focus of today's lecture)**

Behavioral cloning

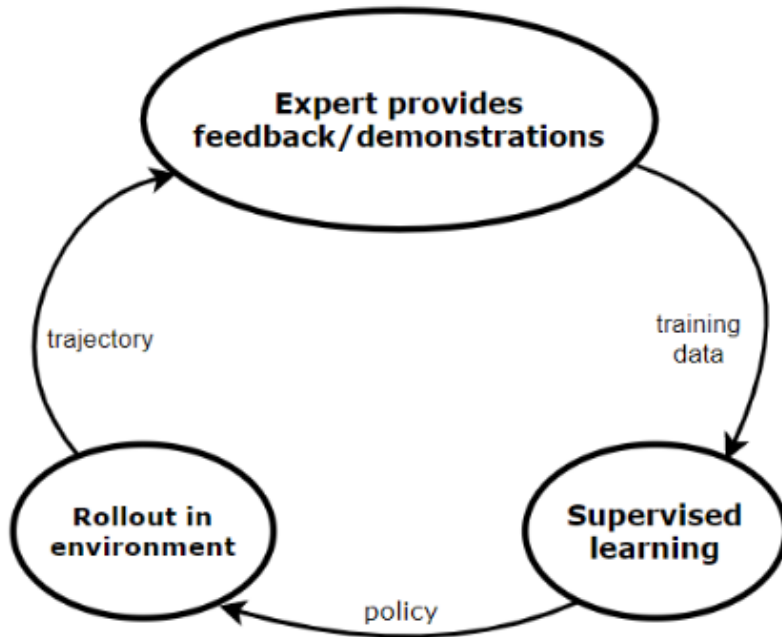
- Simple and efficient form of imitation learning
- Learn the expert policy directly with supervised learning

<https://youtu.be/H0igiP6Hg1k?t=464>

- Problem of “brittleness”: 1-step deviations can lead to huge errors
- Behavioral cloning will at best duplicate the expert’s performance, not exceed it



Direct policy learning (via interactive demonstrator)





- Does not suffer from the same problems as behavioural cloning
- But requires an interactive demonstrator/expert at all times
- Still not clear *why* the agent should perform any given action

Inverse Reinforcement Learning (IRL)

- What if, instead of learning the policy, we learn the reward?

IRL: An informal definition

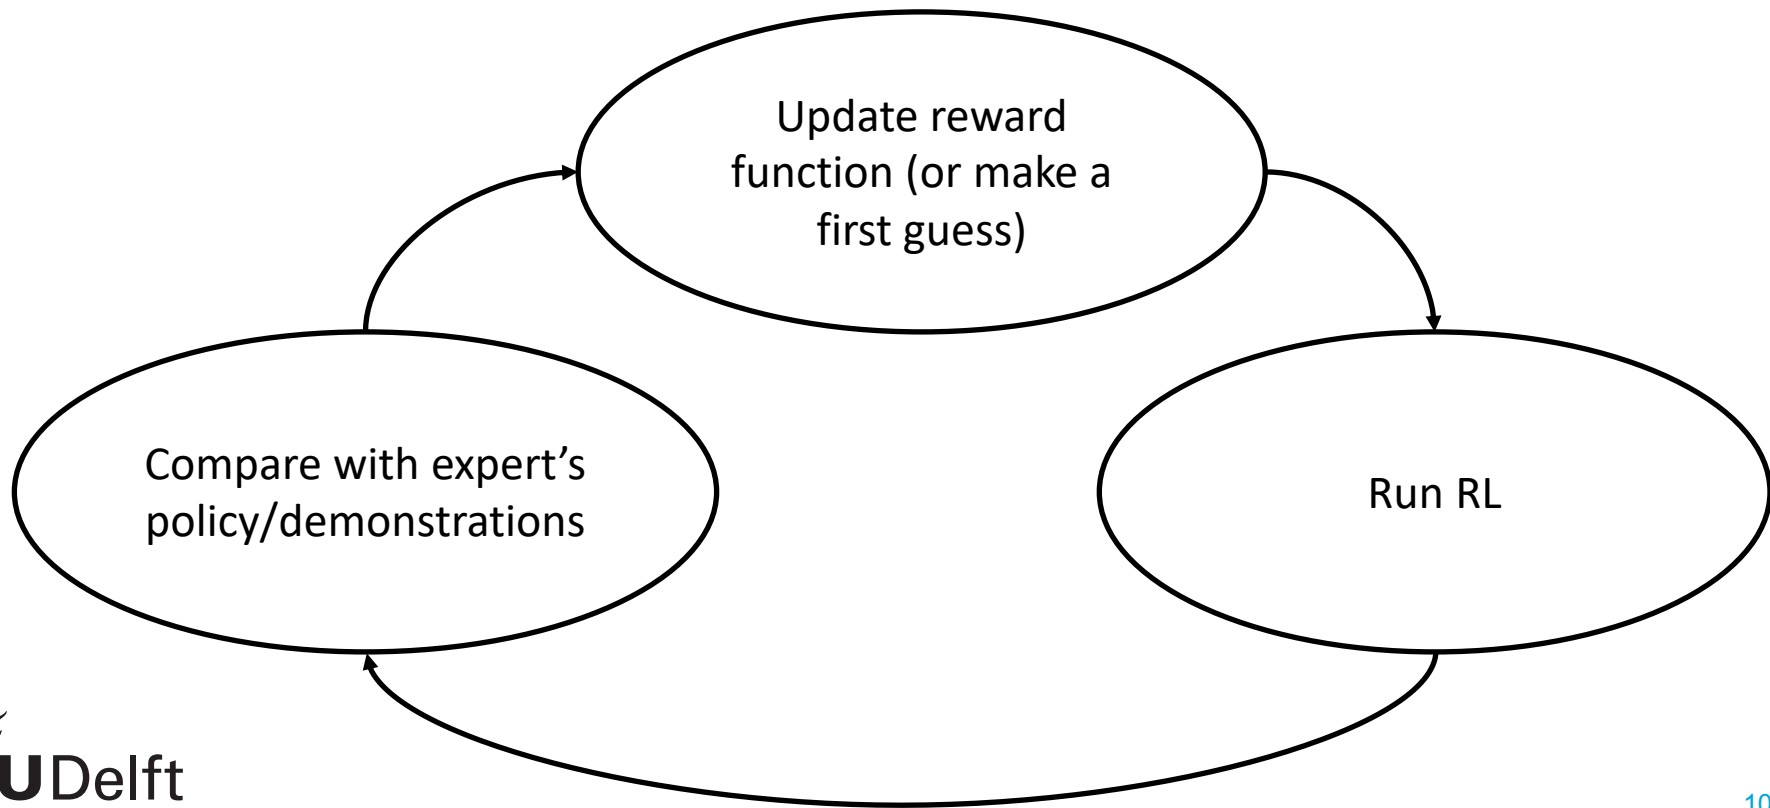


	Direct Policy Learning	Reward Learning	Access to Environment	Interactive Demonstrator	Pre-collected demonstrations
Behavioral cloning (BC)	Yes	No	No	No	Yes
Direct policy learning (interactive IL)	Yes	No	Yes	Yes	Optional
Inverse Reinforcement Learning (IRL)	No	Yes	Yes	No	Yes
Preference-based RL	No	Yes	Yes	Yes	No

Source: Adapted from ICML 2018 Imitational Learning Tutorial. Yisong Yue, Hoang M. Le
https://drive.google.com/file/d/12QdNmMII-bGISWnm8pmd_TawuRN7xagX/viewv
<https://sites.google.com/view/icml2018-imitation-learning/>



* For more information about Preference-based RL, more specifically about the technique reinforcement learning from human feedback RLFH, see: <https://arxiv.org/pdf/1706.03741.pdf>

IRL: An informal definition



IRL Gridworld example

Let's assume that: "Experts" achieve identical or higher rewards than others



Gridworld example

First guess:

- White = 0
- Blue = 1
- Green = 3

Route 1: $0 + 0 + 1 + 3 = 4$

Route 2: $0 + 1 + 1 + 3 = 5$




Second guess:

- White = 0
- Blue = -1
- Green = 2

Route 1: $0 + 0 - 1 + 2 = 1$

Route 2: $0 - 1 - 1 + 2 = 0$



Gridworld example

Third guess:

- White = 0
- Blue = 0
- Green = 1

Route 1: $0 + 0 + 0 + 1 = 1$

Route 2: $0 + 0 + 0 + 1 = 1$




Fourth guess:

- White = 0
- Blue = 0
- Green = 0


Route 1: $0 + 0 + 0 + 0 = 0$

Route 2: $0 + 0 + 0 + 0 = 0$



IRL

- Goal:
 - Find R where π provided by the expert is optimal (let us start with that.. more advanced methods consider demonstrations instead of the policy π)
- But, this is an underdetermined problem. We need some heuristics:
 - Prefer solutions where the expert policy has the largest difference to the other ones
$$\max(value^* - value^{2nd\ best})$$
 - Prefer solutions with smaller rewards
$$\min Reward$$
$$\max(-Reward)$$



Formalizing

- Bellman equation:
$$V^\pi(s) = R(s) + \gamma \sum P_{s\pi(s)}(s') V^\pi(s')$$

- Given that $\pi(s) \equiv a$

- We can rewrite the equation above as:

$$V^\pi = R + \gamma P_{a^*} V^\pi$$

$$V^\pi - \gamma P_{a^*} V^\pi = R$$

$$V^\pi (I - \gamma P_{a^*}) = R$$

$$V^\pi = (I - \gamma P_{a^*})^{-1} R$$

Where:

P_{a^*} is the transition probability matrix, $N \times N$

V^π and R (reward) are $N \times 1$ vectors

γ is the discount factor

Note: In the previous lecture T_a was used for the probability transition matrix. Other notations can also be slightly different. In this and the following slides I use the notation from Ng and Russel (2000)

Formalizing

- Now let's formalize our assumption that π^* achieves identical or higher expected value than all other policies:

$$P_{a^*} V^\pi \geq P_a V^\pi, \forall a \in A \setminus a^*$$

$$P_{a^*} V^\pi - P_a V^\pi \geq 0, \forall a \in A \setminus a^*$$

$$P_{a^*} (I - \gamma P_{a^*})^{-1} R - P_a (I - \gamma P_{a^*})^{-1} R \geq 0, \forall a \in A \setminus a^*$$

$$(P_{a^*} - P_a) (I - \gamma P_{a^*})^{-1} R \geq 0, \forall a \in A \setminus a^*$$

Formalizing

Heuristics:

- Prefer solutions where the expert policy performs better than the other ones
 - Maximize the gap of expected value of acting optimally and the best expected value acting suboptimally

$$\text{maximize} \sum_{i=1}^N \min_{a \in A \setminus a^*} (\mathbf{P}_{a^*} - \mathbf{P}_a) (\mathbf{I} - \gamma \mathbf{P}_{a^*})^{-1} \mathbf{R}$$

- Prefer solutions with smaller rewards
 - Add a penalty term

$$\text{maximize} \sum_{i=1}^N \min_{a \in A \setminus a^*} \{ (\mathbf{P}_{a^*} - \mathbf{P}_a) (\mathbf{I} - \gamma \mathbf{P}_{a^*})^{-1} \mathbf{R} \} - \lambda \|\mathbf{R}\|_1$$

Formal definition

Linear programming formulation:



$$\text{maximize } \sum_{i=1}^N \min_{a \in A \setminus a^*} (\mathbf{P}_{a^*} - \mathbf{P}_a) (\mathbf{I} - \gamma \mathbf{P}_{a^*})^{-1} \mathbf{R} - \lambda \|\mathbf{R}\|_1$$

$$\begin{aligned} \text{s. t. } & (\mathbf{P}_{a^*} - \mathbf{P}_a) (\mathbf{I} - \gamma \mathbf{P}_{a^*})^{-1} \mathbf{R} \geq 0, \forall a \in A \setminus a^* \\ & |R_i| \leq R_{\max}, i = 1, \dots, N \end{aligned}$$

A practical example

$$\begin{aligned} & \underset{\pi}{\text{maximize}} \sum_{i=1}^N \min_{a \in A \setminus a^*} (\mathbf{P}_{a^*} - \mathbf{P}_a) (\mathbf{I} - \gamma \mathbf{P}_{a^*})^{-1} \mathbf{R} - \lambda \|\mathbf{R}\|_1 \\ & \text{s.t. } (\mathbf{P}_{a^*} - \mathbf{P}_a) (\mathbf{I} - \gamma \mathbf{P}_{a^*})^{-1} \mathbf{R} \geq 0, \forall a \in A \setminus a^* \\ & \quad |R_i| \leq R_{\max}, i = 1, \dots, N \end{aligned}$$

5x5 gridworld environment, where π^* :




Actions: (, , , )

- Agents start from the lower-left grid square (0,0), and finish on the upper-right grid square (4,4)
- Agents move (up, down, left, right) just one square at a time
- Actions have a 30% chance of moving in a random direction (wind)
- Discount factor $\gamma = 0.2$
- Penalty factor $\lambda = 1.05$
- $R_{\max} = 1$


A practical example

$$\begin{aligned} & \underset{P_a}{\text{maximize}} \sum_{i=1}^N \min_{a \in A \setminus a^*} (P_{a^*} - P_a) (I - \gamma P_{a^*})^{-1} R - \lambda \|R\|_1 \\ & \text{s.t. } (P_{a^*} - P_a) (I - \gamma P_{a^*})^{-1} R \geq 0, \forall a \in A \setminus a^* \\ & |R_i| \leq R_{\max}, i = 1, \dots, N \end{aligned}$$

5x5 gridworld environment, where π^* :



Actions: (↑, →, ↓, ←)



- Step 1: Calculate P_a

$$P_{a=\uparrow}(0,0)(0,0) = 0$$


$$P_{a=\uparrow}(0,0)(0,1) = 1 - 0.3 + 0.3/4$$

$$P_{a=\uparrow}(0,0)(1,0) = 0.3/4$$

$$P_{a=\uparrow}(0,0)(0,2) = 0$$

....

- Step 2: Run the linear IRL algorithm



Exercise sheet on Brightspace:

A working example of IRL is provided in the tutorial sheet


Explore the code and answer the questions

- This linear formulation requires an optimal policy, but we usually all we have access is a set of trajectories (demonstrations)
- Can we do better than the heuristics we used?

Maximum Entropy IRL

- Handle ambiguity using a probabilistic model of behavior
- Employs the *principle of maximum entropy* to resolve the ambiguity in choosing a distribution over decisions in a principled way

Principle of maximum entropy: the probability distribution which best represents the current state of knowledge about a system is the one with largest entropy.



Maximum Entropy IRL

- Notion of suboptimality
 - If the demonstrator has some random behavior at times, this might mean that they don't care much about specific actions in this setting → Lower/No reward
 - If the demonstrator consistently does a specific action in a given setting, it probably means that they really care about it → Larger reward
- The idea is to match the most relevant features, but besides that be as random as possible

Maximum Entropy IRL

- A bit more formal.. We want to learn a reward function that yields the distribution over all trajectories with the maximum entropy:


$$\max_D - \sum_{\pi \in (S \times A)} Pr(\pi) \log Pr(\pi)$$

- While respecting two constraints:
 - The distribution over all trajectories should be a probability distribution
 - The expected feature count of the demonstrated trajectories must match the empirical feature

- Maximum Entropy IRL deals with the ambiguity problem in a more principled manner
- But, still it only consider linear reward functions
- Other approaches can deal with non-linear reward functions, for example...

Adversarial IRL

- Trains a policy against a discriminator that aims to distinguish the expert demonstrations from the learned policy
- Adversarial IRL can deal with non-linear rewards



- Finally, this process could be done in a more collaborative manner
- Humans teach each other all the time... and someone that is teaching something will not necessarily behave “optimally”, but will try to modify the behavior to support learning



Collaborative IRL

- Cooperative, partial information game with two agents, human and robot
 - Human knows the reward function
 - Robot does not know the reward function
- The human and the robot get rewards determined by the same reward function → incentivizes the human to teach and the robot to learn



Critiques and limitations

- Underdetermined problem (ambiguity)
 - Wrong guesses can lead to high regret
- Difficult to evaluate a learned reward \rightarrow we do not have access to a “ground truth reward”
- Demonstrations may not be optimal



Critiques and limitations

- IRL assumes that human behavior is optimal or noisily optimal
- However, humans often deviate from such rationality assumptions:
 - In systematic, non-random ways: Biases such as time inconsistency, loss aversion and anchoring
 - But also due to cognitive aspects such as forgetfulness, limited planning and false beliefs



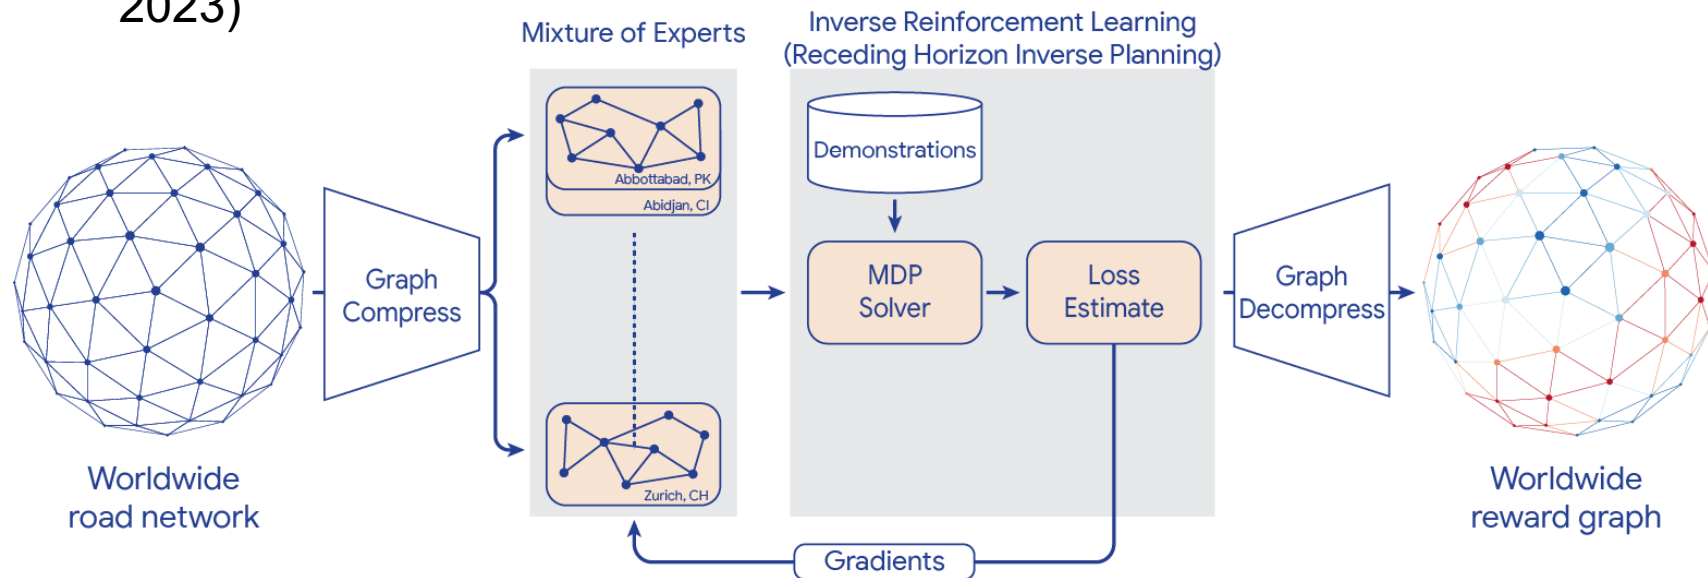
Critiques and limitations

- Methods have been proposed to account for specific human deviations from such rational expectations e.g. specific biases or noise rationality, but no general framework has been proposed
- Mindermann and Armstrong (2018) demonstrated that:
 - It is impossible to uniquely decompose a policy (or demonstrations) into a planning algorithm (i.e. human thinking/"rationality") and a reward function
 - Normative assumptions are needed, which cannot be deduced exclusively from observations



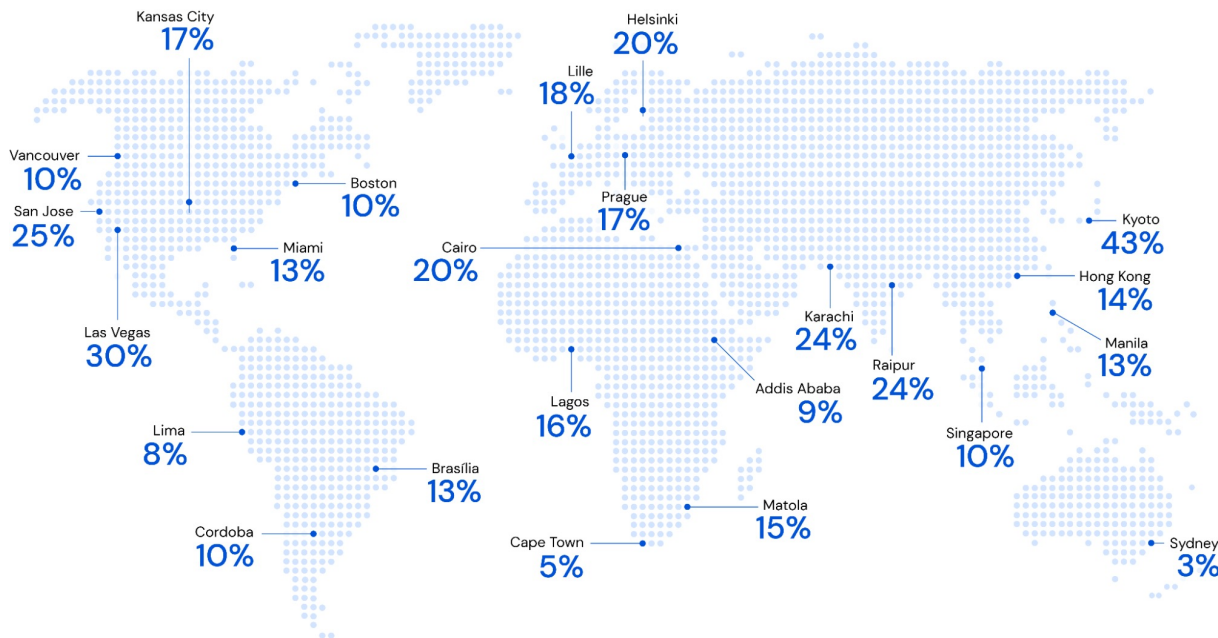
A recent real-world application...

- World scale inverse reinforcement learning in Google Maps (September 12, 2023)



A recent real-world application...

- World scale inverse reinforcement learning in Google Maps (September 12, 2023)



Wrapping up



- IRL methods can learn a reward function from human demonstrations
- However, especially in social complex environments, it is important to take very well into consideration that such algorithms might not model all nuances and particularities of human behavior -> Model \neq Reality

References

IRL:

- A. Y. Ng and S. J. Russell. 2000. Algorithms for inverse reinforcement learning. In: *Proceedings of the 17th International Conference on Machine Learning (ICML '00)*, Stanford University, Stanford, CA, USA.
<https://ai.stanford.edu/~ang/papers/icml00-irl.pdf>
- Ziebart, B. D., Maas, A. L., Bagnell, J. A., & Dey, A. K. (2008, July). Maximum entropy inverse reinforcement learning. In *Aaai* (Vol. 8, pp. 1433-1438).
- Arora, S., & Doshi, P. (2021). A survey of inverse reinforcement learning: Challenges, methods and progress. *Artificial Intelligence*, 297, 103500.

RL text book:

- R. S. Sutton and A. G. Barto. 2018. Reinforcement learning: An introduction (2nd ed). MIT press.
<https://www.andrew.cmu.edu/course/10-703/textbook/BartoSutton.pdf>

Text book on linear programming:

- R. J. Vanderbei. 2020. Linear programming: foundations and extensions (5th ed). Springer Nature.

Inverse Reinforcement Learning

CS4375 Artificial Intelligence Techniques

Luciano Cavalcante Siebert

Assistant Professor

Interactive Intelligence (II) Group, INSY, EEMCS, TU Delft

l.cavalcantesiebert@tudelft.nl